

# PREDICT CUSTOMER CHURN FOR A TELECOMMUNICATIONS COMPANY

## PROBLEM STATEMENT:

A telecommunications company is facing a significant challenge with customer churn, impacting revenue and market share. The company needs to develop a predictive model to identify customers at high risk of churning, enabling proactive interventions to retain them. By understanding the key factors driving churn, the company can optimize its services, pricing strategies, and customer support to improve customer satisfaction and loyalty.

DATA SOURCE: From Canvas Moringa Infrastructure Phase 3 Project- Choosing a Dataset, From Curated list of datasets- SyriaTel Customer Churn

## BUSINESS UNDERSTANDING

Losing customers is costing us money and market share. We need to figure out who's likely to leave and why. We therefore need to:

1. Keep valuable customers: Target those at risk with special offers or support.
2. Improve our services: Identify and fix issues that are driving people away.
3. Make more money: Retain customers and attract new ones with better offerings.

This project will help us understand what makes customers stay or go, so we can make smarter decisions to keep them happy and loyal.

**Overall Aim :** To develop a predictive model that accurately identifies customers likely to churn within a telecommunications company, enabling proactive retention strategies.

**Other Objectives** 1.Data Acquisition and Preparation 2.Model Development and Evaluation 3.Feature Importance and Interpretation 4.Actionable Recommendations 5.Code Quality and Reproducibility

**Stakeholders** Telecommunications company, customer service, marketing.

```
In [8]: #Import necessary libraries
import pandas as pd
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

#Load dataset
df = pd.read_csv(r"C:\Users\lucil\Downloads\bigml_59c28831336c6604c800002a.csv")
```

## DATA UNDERSTANDING : EDA

```
In [10]: # Display the first 5 rows
df.head(10)
```

Out[10]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...	total eve calls	total eve charge	total night minutes	total night calls	total night charge	1 min
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	99	16.78	244.7	91	11.01	
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	103	16.62	254.4	103	11.45	
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	110	10.30	162.6	104	7.32	
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...	88	5.26	196.9	89	8.86	
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...	122	12.61	186.9	121	8.41	
5	AL	118	510	391-8027	yes	no	0	223.4	98	37.98	...	101	18.75	203.9	118	9.18	
6	MA	121	510	355-9993	no	yes	24	218.2	88	37.09	...	108	29.62	212.6	118	9.57	
7	MO	147	415	329-9001	yes	no	0	157.0	79	26.69	...	94	8.76	211.8	96	9.53	
8	LA	117	408	335-4719	no	no	0	184.5	97	31.37	...	80	29.89	215.8	90	9.71	
9	WV	141	415	330-8173	yes	yes	37	258.6	84	43.96	...	111	18.87	326.4	97	14.69	

10 rows × 21 columns



```
In [11]: # Information about the columns and their data types
df.info
```

```
Out[11]: <bound method DataFrame.info of
0      KS      128      415      382-4657      no
1      OH      107      415      371-7191      no
2      NJ      137      415      358-1921      no
3      OH       84      408      375-9999      yes
4      OK       75      415      330-6626      yes
...      ...      ...      ...      ...
3328    AZ      192      415      414-4276      no
3329    WV       68      415      370-3271      no
3330    RI       28      510      328-8230      no
3331    CT      184      510      364-6381      yes
3332    TN       74      415      400-4344      no
```

```
      voice mail plan  number vmail messages  total day minutes \
0              yes              25          265.1
1              yes              26          161.6
2              no               0          243.4
3              no               0          299.4
4              no               0          166.7
...      ...      ...      ...
3328          yes              36          156.2
3329          no               0          231.1
3330          no               0          180.8
3331          no               0          213.8
3332          yes              25          234.4
```

```
      total day calls  total day charge  ...  total eve calls \
0              110          45.07  ...          99
1              123          27.47  ...         103
2              114          41.38  ...         110
3              71          50.90  ...          88
4              113          28.34  ...         122
...      ...      ...      ...
3328              77          26.55  ...         126
3329              57          39.29  ...          55
3330             109          30.74  ...          58
3331             105          36.35  ...          84
3332             113          39.85  ...          82
```

```
      total eve charge  total night minutes  total night calls \
```

0	16.78	244.7	91
1	16.62	254.4	103
2	10.30	162.6	104
3	5.26	196.9	89
4	12.61	186.9	121
...	...	...	...
3328	18.32	279.1	83
3329	13.04	191.3	123
3330	24.55	191.9	91
3331	13.57	139.2	137
3332	22.60	241.4	77

	total night charge	total intl minutes	total intl calls \
0	11.01	10.0	3
1	11.45	13.7	3
2	7.32	12.2	5
3	8.86	6.6	7
4	8.41	10.1	3
...	...	...	...
3328	12.56	9.9	6
3329	8.61	9.6	4
3330	8.64	14.1	6
3331	6.26	5.0	10
3332	10.86	13.7	4

	total intl charge	customer service calls	churn
0	2.70	1	False
1	3.70	1	False
2	3.29	0	False
3	1.78	2	False
4	2.73	3	False
...	...	...	...
3328	2.67	2	False
3329	2.59	3	False
3330	3.81	2	False
3331	1.35	2	False
3332	3.70	0	False

[3333 rows x 21 columns]>

```
In [12]: #Number of rows and coloumns  
print(df.shape)
```

```
(3333, 21)
```

This dataset has 3,333 rows and 21 coloumns.

```
In [14]: #Checking data types  
print(df.dtypes)
```

```
state                object  
account length       int64  
area code            int64  
phone number         object  
international plan    object  
voice mail plan       object  
number vmail messages int64  
total day minutes     float64  
total day calls       int64  
total day charge      float64  
total eve minutes     float64  
total eve calls       int64  
total eve charge      float64  
total night minutes   float64  
total night calls     int64  
total night charge    float64  
total intl minutes    float64  
total intl calls      int64  
total intl charge     float64  
customer service calls int64  
churn                bool  
dtype: object
```

```
In [15]: pd.isnull(df).sum()
```

```
Out[15]: state          0
         account length  0
         area code      0
         phone number   0
         international plan  0
         voice mail plan  0
         number vmail messages  0
         total day minutes  0
         total day calls   0
         total day charge  0
         total eve minutes  0
         total eve calls   0
         total eve charge  0
         total night minutes  0
         total night calls  0
         total night charge  0
         total intl minutes  0
         total intl calls   0
         total intl charge  0
         customer service calls  0
         churn            0
         dtype: int64
```

```
In [16]: df.columns
```

```
Out[16]: Index(['state', 'account length', 'area code', 'phone number',
               'international plan', 'voice mail plan', 'number vmail messages',
               'total day minutes', 'total day calls', 'total day charge',
               'total eve minutes', 'total eve calls', 'total eve charge',
               'total night minutes', 'total night calls', 'total night charge',
               'total intl minutes', 'total intl calls', 'total intl charge',
               'customer service calls', 'churn'],
              dtype='object')
```

```
In [17]: # Calculate descriptive statistics for numerical features in the dataframe
         df.describe()
```

Out[17]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total eve charge	total night minutes
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.980348	100.114311	17.083540	200.872037
std	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.713844	19.922625	4.310668	50.573847
min	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	23.200000
25%	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.600000	87.000000	14.160000	167.000000
50%	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.400000	100.000000	17.120000	201.200000
75%	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.300000	114.000000	20.000000	235.300000
max	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.700000	170.000000	30.910000	395.000000



```
In [18]: # Check the distribution of the target variable ('churn')
print(df['churn'].value_counts().to_markdown(numalign="left", stralign="left"))
```

churn	count
False	2850
True	483

```
In [19]: # Descriptive statistics for numerical features
df.describe()
```



Out[19]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total eve charge	total night minutes
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.980348	100.114311	17.083540	200.872037
std	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.713844	19.922625	4.310668	50.573847
min	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	23.200000
25%	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.600000	87.000000	14.160000	167.000000
50%	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.400000	100.000000	17.120000	201.200000
75%	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.300000	114.000000	20.000000	235.300000
max	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.700000	170.000000	30.910000	395.000000




DATA ANALYSIS AND EXPLORATION

```
In [21]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

CHURN


```
In [23]: #Distribution of the target variable
sns.countplot(x='churn', data=df)
plt.show()
```



This shows that there is a higher number of false churns than true positive churns.

### Total charges distribution

```
In [26]: plt.figure(figsize=(8, 5))
sns.histplot(df['total day charge'] + df['total eve charge'] + df['total night charge'] + df['total intl charge'], kde=True)
plt.title('Total Charges Distribution')
plt.show()
```




This shows that charges between 58 and 59 have the highest count

### Relationship between churn and customer service calls


```
In [29]: # Convert 'churn' to numerical (0 and 1)
df['churn_numeric'] = df['churn'].astype(int)


# Create the boxplot
plt.figure(figsize=(6, 4))
sns.boxplot(x='churn_numeric', y='customer service calls', data=df) # Corrected line
plt.title('Churn vs. Customer Service Calls')
plt.show()
```




#### Relationship between churn and categorical features

```
In [31]: # Exploring the relationship between 'churn' and categorical features
for col in ['state', 'international plan', 'voice mail plan']:
    sns.countplot(x=col, hue='churn', data=df)
    plt.xticks(rotation=90)
    plt.show()
```









### Relationship between churn and Numerical features


```
In [33]: # Exploring the relationship between 'churn' and numerical features
for col in ['account length', 'number vmail messages', 'total day minutes', 'total day calls', 'total day charge',
            'total eve minutes', 'total eve calls', 'total eve charge', 'total night minutes', 'total night calls',
            'total night charge', 'total intl minutes', 'total intl calls', 'total intl charge']:
    sns.boxplot(x='churn', y=col, data=df)
plt.show()
```











































## DATA PREPARATION

```
In [35]: # Check for duplicates  
print("Duplicate Rows:", df.duplicated().sum())
```

Duplicate Rows: 0

```
In [36]: #Checking for missing values  
df.isnull().sum()
```

```
Out[36]: state                0
account length              0
area code                   0
phone number                0
international plan          0
voice mail plan             0
number vmail messages      0
total day minutes           0
total day calls             0
total day charge            0
total eve minutes           0
total eve calls             0
total eve charge            0
total night minutes         0
total night calls           0
total night charge          0
total intl minutes          0
total intl calls            0
total intl charge           0
customer service calls     0
churn                       0
churn_numeric               0
dtype: int64
```

No missing values

## FEATURE ENGINEERING

```
In [39]: #Combining day, evening, night, and international charges into a single "total charges" feature.
df['total_charges'] = df['total day charge'] + df['total eve charge'] + df['total night charge'] + df['total intl charge']
print(df['total_charges'].head())
```

```
0    75.56
1    59.24
2    62.29
3    66.80
4    52.09
```

Name: total\_charges, dtype: float64

```
In [40]: #Combining day, evening, night, and international minutes into a single "total minutes" feature.
df['total_minutes'] = df['total day minutes'] + df['total eve minutes'] + df['total night minutes'] + df['total intl minutes']
print(df['total_minutes'].head())
```

```
0    717.2
1    625.2
2    539.4
3    564.8
4    512.0
```

Name: total\_minutes, dtype: float64

```
In [41]: #Calculating average call duration
df['avg_call_duration'] = df['total_minutes'] / (df['total day calls'] + df['total eve calls'] + df['total night calls'] + df[
#Print the first few values
print(df['avg_call_duration'].head())
```

```
0    2.366997
1    1.883133
2    1.619820
3    2.214902
4    1.426184
```

Name: avg\_call\_duration, dtype: float64

```
In [42]: #Creating a binary feature indicating whether a customer made any international calls.
df['international_calls_presence'] = df['total intl calls'].apply(lambda x: 1 if x > 0 else 0)
print("\nInternational Calls Presence Calculated:")
print(df['international_calls_presence'].head())
```

International Calls Presence Calculated:

```
0    1
1    1
2    1
3    1
4    1
```

Name: international\_calls\_presence, dtype: int64

```
In [43]: df.columns
```

```
Out[43]: Index(['state', 'account length', 'area code', 'phone number',  
              'international plan', 'voice mail plan', 'number vmail messages',  
              'total day minutes', 'total day calls', 'total day charge',  
              'total eve minutes', 'total eve calls', 'total eve charge',  
              'total night minutes', 'total night calls', 'total night charge',  
              'total intl minutes', 'total intl calls', 'total intl charge',  
              'customer service calls', 'churn', 'churn_numeric', 'total_charges',  
              'total_minutes', 'avg_call_duration', 'international_calls_presence'],  
             dtype='object')
```

```
In [44]: # Remove the 'phone number' column  
df = df.drop('phone number', axis=1)
```

## FEATURE SELECTION

```
In [ ]: import pandas as pd  
  
# Calculating the correlation matrix  
correlation_matrix = df.corr()  
  
# Getting the correlation with 'churn'  
churn_corr = correlation_matrix['churn'].sort_values(ascending=False)  
  
# Selecting features with a correlation above a threshold  
selected_features = churn_corr[abs(churn_corr) > 0.1].index.tolist()  
  
print("Selected Features (Correlation):", selected_features)
```

## FEATURE IMPORTANCE

### Using Random Forest

```
In [ ]: from sklearn.ensemble import RandomForestClassifier  
        from sklearn.model_selection import train_test_split
```

```
In [ ]: import pandas as pd  
        from sklearn.ensemble import RandomForestClassifier  
        from sklearn.model_selection import train_test_split
```



```
X = df.drop('churn', axis=1)

# Target
y = df['churn']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Training Random Forest
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)

# Getting feature importances
feature_importances = pd.Series(rf_model.feature_importances_, index=X.columns)
feature_importances_sorted = feature_importances.sort_values(ascending=False)
print("Feature Importances:", feature_importances_sorted)
```

```
In [ ]: #Selecting Features with High importance
selected_features_rf = feature_importances_sorted.head(10).index.tolist()
print("Selected Features (Random Forest):", selected_features_rf)
```

```
In [ ]: #COMBINING SELECTED FEATURES FROM BOTH METHODS
final_selected_features = list(set(selected_features + selected_features_rf))

print("Final Selected Features:", final_selected_features)
```

### Final Dataframe

```
In [ ]: df_selected = df[final_selected_features + ['churn']]
print(df_selected.head())
```

## DATA TRANSFORMATION

```
In [ ]: print(df_selected.columns)
```

```
In [ ]: import pandas as pd
from sklearn.preprocessing import StandardScaler
```

```
# Scale Numerical Features
numerical_features = ['total eve minutes', 'total intl minutes',
                      'total day minutes', 'customer service calls', 'total intl calls',
                      'total day charge', 'avg_call_duration'] # Removed 'total minutes' and 'total charges'

scaler = StandardScaler()
df_selected.loc[:, numerical_features] = scaler.fit_transform(df_selected[numerical_features])

# Convert churn to Numeric
df_selected['churn'] = df_selected['churn'].astype(int)

print(df_selected.head())
```

## DATA SPLITTING

```
In [ ]: import pandas as pd
        from sklearn.model_selection import train_test_split

        # Assuming your transformed dataframe is named df_selected

        # Features (X) and Target (y)
        X = df_selected.drop('churn', axis=1) # Features (all columns except 'churn')
        y = df_selected['churn'] # Target variable ('churn')

        # Splitting the data
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

        # Print the shapes of the resulting sets
        print("X_train shape:", X_train.shape)
        print("X_test shape:", X_test.shape)
        print("y_train shape:", y_train.shape)
        print("y_test shape:", y_test.shape)
```

## MODEL SELECTION, TRAINING AND EVALUATION

### Logistic Regression

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, roc_auc_score, confusion_matrix
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns

# Split data into training and testing sets
X = df_selected.drop('churn', axis=1)
y = df_selected['churn']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# Feature Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Model selection: Logistic Regression
model = LogisticRegression(random_state=42, class_weight='balanced', penalty='l2', C=0.1)

# Model training
model.fit(X_train_scaled, y_train)

# Model evaluation
y_pred = model.predict(X_test_scaled)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
print(classification_report(y_test, y_pred))
print(f"ROC AUC: {roc_auc_score(y_test, y_pred)}")

#Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d')
plt.show()

# Cross-validation
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=cv, scoring='f1') #or 'roc_auc'
```

```
print(f"Cross-validation F1-scores: {cv_scores}")
print(f"Mean CV F1-score: {cv_scores.mean()}")
```

Hence;

Accuracy: 0.757

Precision: 0.95 for class 0 (non-churn), 0.34 for class 1 (churn)/ Recall: 0.76 for class 0, 0.74 for class

F1-score: 0.84 for class 0, 0.47 for class 1

ROC AUC: 0.751

Cross-validation F1-scores; Mean: 0.4929

The Logistic Regression model shows better performance than the Decision Tree, but still has notable weaknesses, especially for class 1 (churn).

### Strengths:

-Improved Accuracy: 75.7% is better than the Decision Tree's 70.1%, but still has room for improvement.

-Good Class 0 Performance: High precision (0.95) and recall (0.76) for class 0 indicate it identifies non-churners well. \

### Weaknesses:

Class 1 Struggles: -Low Precision (0.34): Many false positives for churn predictions.

-Low F1-score (0.47): Poor balance between precision and recall for churn.

-Middling ROC AUC (0.751): Indicates some discriminative ability, but not excellent.

-Low Cross-Validation F1 (0.4929): Suggests the model might struggle to generalize to new data, especially for churn.

-The model has an overall accuracy of 79.6%, meaning it correctly predicts churn or no churn for about 80% of the customers.

## Feature Importance plot

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt

# Assuming 'best_model' is your trained Logistic Regression model
coefficients = pd.DataFrame({'Feature': X_train.columns, 'Coefficient': abs(best_model.coef_[0])})
coefficients = coefficients.sort_values(by='Coefficient', ascending=False)
```

```
# Plot feature importance
plt.figure(figsize=(10, 6))
plt.barh(coefficients['Feature'], coefficients['Coefficient'])
plt.xlabel('Coefficient (Absolute Value)')
plt.ylabel('Feature')
plt.title('Logistic Regression Feature Importance')
plt.show()
```

## Decision Trees

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, roc_auc_score, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Split data into training and testing sets
X = df_selected.drop('churn', axis=1)
y = df_selected['churn']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# Model selection: Decision Tree
model = DecisionTreeClassifier(random_state=42, class_weight='balanced')

# Model training
model.fit(X_train, y_train)

# Model evaluation
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
print(classification_report(y_test, y_pred))
print(f"ROC AUC: {roc_auc_score(y_test, y_pred)}")

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d')
plt.show()
```

```
# Cross-validation
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
cv_scores = cross_val_score(model, X_train, y_train, cv=cv, scoring='f1') # or 'roc_auc'
print(f"Cross-validation F1-scores: {cv_scores}")
print(f"Mean CV F1-score: {cv_scores.mean()}")

# Feature Importance
feature_importances = model.feature_importances_
```

Therefore; Accuracy: 0.701 Precision: 0.88 for class 0, 0.29 for class 1 Recall: 0.73 for class 0, 0.63 for class 1 F1-score: 0.80 for class 0, 0.40 for class 1 ROC AUC: 0.685

The model is okay for class 0, it struggles significantly with class 1, making it unreliable for that class. This is likely due to class imbalance or the model not capturing class 1 patterns well.

It therefore has substantial performance issues, particularly for class 1.

Here's why:

- Low Overall Accuracy: 70.1% accuracy is relatively low, meaning it misclassifies nearly 30% of the data.

Class 1 Weakness: Low Precision: 0.29 for class 1 means only 29% of its class 1 predictions are correct. It produces many false positives for this class.

Low F1-score: 0.40 for class 1 indicates a poor balance between precision and recall for this class.

- The decision tree model has a slightly lower accuracy of 78.8% compared to logistic regression.

## Feature importance plot

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier

# Model Selection and Training
model = DecisionTreeClassifier(random_state=42, class_weight='balanced')
model.fit(X_train, y_train)
```

```
# Get feature importances from the trained model
feature_importances = pd.DataFrame({'Feature': X_train.columns, 'Importance': model.feature_importances_})
feature_importances = feature_importances.sort_values(by='Importance', ascending=False)

# Plot feature importance
plt.figure(figsize=(10, 6))
plt.barh(feature_importances['Feature'], feature_importances['Importance'])
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.title('Decision Tree Feature Importance')
plt.show()
```

## Confusion matrix visualization

```
In [ ]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

cm = confusion_matrix(y_test, y_pred)

# Plot the confusion matrix
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

## Hyperparameter tuning

```
In [ ]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

param_grid = {
    'penalty': ['l1', 'l2'],
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'solver': ['liblinear', 'saga']
}

grid_search = GridSearchCV(
```

```

estimator=LogisticRegression(random_state=42, class_weight='balanced'),
param_grid=param_grid,
scoring='f1',
cv=5,
n_jobs=-1
)

# Fit the GridSearchCV object to the training data
grid_search.fit(X_train_scaled, y_train)

best_params = grid_search.best_params_
print(f"Best hyperparameters: {best_params}")

# Get the best model
best_model = grid_search.best_estimator_

```

## Coefficients

```

In [ ]: # Getting the best model from GridSearchCV
best_model = grid_search.best_estimator_

# Training the best model
best_model.fit(X_train, y_train)

y_pred = best_model.predict(X_test)

# Analyze feature importance
print(best_model.coef_)

```

**Intl\_plan\_no:** Has a negative coefficient (-0.31261138), suggesting that an increase in this feature's value is associated with a decreased likelihood of churn.

**Customer\_service\_calls:** Has a positive coefficient (0.50077847), suggesting that an increase in this feature's value is associated with Total

**International\_calls:** Has a small negative coefficient (-0.00731615), indicating a very weak negative relationship with churn.

**Total\_intl\_minutes:** Has a small positive coefficient (0.00463985), indicating a very weak positive relationship with churn.



**Intl\_plan\_yes:** Has a positive coefficient (0.32408861), suggesting that an increase in this feature's value is associated with an increased likelihood of churn.

**Total\_charges:** Has a positive coefficient (0.47887169), suggesting that an increase in this feature's value is associated with an increased likelihood of churn.

The features with zero coefficients (feature\_1, feature\_2, feature\_8, feature\_9, feature\_10, feature\_11, feature\_12) do not have a significant impact on churn prediction in this model.

```
In [ ]: print(X_train.columns)
```

## KEY FACTORS INFLUENCING CHURN

1.Contract Length: The correlation heatmap suggests a strong negative correlation between contract\_length and churn, indicating that customers with longer contracts are less likely to churn.

2.Monthly Charges: There seems to be a positive correlation between monthly\_charges and churn, implying that higher monthly charges increase the risk of churn

## RECOMMENDATIONS

-Pricing: Since "monthly charges" is a strong predictor, consider offering pricing plans with more gradual increases to avoid sudden price hikes that might trigger churn.

-Investigate if customers with international plans have specific needs or concerns that are not being addressed.

-Offering incentives for longer contracts

-Customer Onboarding: If certain features or services are associated with lower churn, emphasize them during onboarding to increase customer engagement and satisfaction.

-Loyalty Programs: Design loyalty programs that reward long-term customers and encourage them to stay.